

Project 0

HPC Primer in C++

Due date: See iCorsi submission

This project will introduce you to basic programming concepts for HPC. This is meant to be a short project to harmonize the basic programming skills of incoming students. For this reason this is the only project which does not require a project report.

Please note: This project must be completed individually by each student. All student submissions will be run through a plagiarism checker to ensure that students have done the assignment independently. Please carefully read the “Additional notes and submission details” at the end of this document for instructions on how to submit your work.

1 Compiling and running a program on the Rosa cluster[20 points]

The file `read_environment.cpp` is a simple C++ program that reads some environment variables and prints it to the console. Your first task is to log onto the Rosa cluster, compile, and then run this program. Please submit a screenshot of your terminal after the program has finished execution. To log onto the Rosa cluster you will need to use the [Secure Shell \(SSH\)](#) cryptographic network protocol. This allows you to open up a terminal on a remote machine. You may need to be logged into the [VPN](#) if you are not connected to the USI network.

```
[user@localhost]$ ssh userid@rosa.usi.ch  
[user@icslogin01]$
```

Once you successfully log into the cluster you will be connected to the login node. The login node manages basic administration of the cluster and is the common interface for all users. Running programs on the login node may cause it to bog down and crash, which could bring down the entire cluster. Thus, **you should never run programs on the login node**. You will know that you are connected to the login node if the hostname in your cursor contains the word “login”; for example, in the code snippet shown above the hostname is `icslogin01`.

To connect to a compute node you may use:

```
[user@icslogin01]$ srun --nodes=1 --time=01:00:00 --pty bash -i  
[user@icsnodeXX]$
```

or during the lectures/lab hours you might use our course reservation¹ allocation on Monday 11:00-12:30

```
[user@icslogin01]$ srun --nodes=1 --time=01:00:00 --reservation=hpc-monday --pty bash -i  
[user@icsnodeXX]$
```

and on Thursday 11:00-12:30

```
[user@icslogin01]$ srun --nodes=1 --time=01:00:00 --reservation=hpc-thursday --pty bash -i  
[user@icsnodeXX]$
```

Now that you are connected to a compute node, you must load the GNU compiler to compile your program. To do this please use:

```
[user@icsnodeXX]$ module load gcc
```

Please note that you will need to re-load your required modules each time you connect to a compute node. You can see the modules that are currently loaded using the command:

```
[user@icsnodeXX]$ module list
```

¹Reservations allocate a good portion of the cluster, not using the reservation during the time it is enabled will probably result in longer waiting times

Next you need to compile the program. To do this you will use the GNU C++ compiler `g++` (if we were compiling C code we would use `gcc`).

```
[user@icsnodeXX]$ g++ -o ciao read_environment.cpp
```

What we have done with this command is to run the program “`g++`” and we have passed it three arguments separated by space:

1. The option flag `-o` which notifies the program that the next argument will contain the desired name of the compiled output file.
2. The output filename that we wish to specify.
3. The name of the file that we wish to compile.

You should now have a binary file named `ciao` in your directory. Execute this binary, passing it your first and last name in camelCase as a command line argument.

```
[user@icsnodeXX]$ ./ciao firstNameLastName
```

Take a screenshot of your terminal showing the results and save it as the deliverable for this first problem.

2 Implementing a basic fraction toolbox in C++ [50 points]

In the `fraction_summing` directory you will find the skeleton of a program to do some basic operations with fractions. Header files help you organize your codebase more effectively by moving function and variable declarations to a file separate from the main program. In the header file `fraction_toolbox.hpp` you will find a number of function declarations. In the C/C++ programming language function declarations are generally of the form:

```
functionReturnType functionName(inputDataType inputName);
```

In the directory you will also find the source file `fraction_toolbox.cpp` where the functions are actually defined. Please complete the following tasks in the fraction toolbox files in order to implement a basic fraction toolbox (each item will have a TODO comment in the source code) [5 points each]:

1. In the header file define a C datatype `struct fraction` to store a rational number (i.e., a fraction of integers). The numerator and the denominator of the fraction should be named `num` and `denom` respectively.
2. Write a function `square_fraction()` that takes a `fraction` structure as input and returns its square as a `fraction` while leaving the original input unchanged.
3. Write a function, `square_fraction_inplace()`, that takes a `fraction` structure as input and squares that `fraction` without returning any data. Add the function declaration to `fraction_toolbox.hpp`.
4. Write a function, `fraction2double()`, that takes a `fraction` as input and returns its value as a double precision floating point number.
5. Write a function, `gcd()`, that uses the *recursive* Euclid’s algorithm (see Algorithm 1 below) to return the greatest common divisor of two integers.
6. Write a function, `gcd()` ([overloading](#) the `gcd()` function call), that takes a `fraction` as input and returns the greatest common divisor (integer) of the numerator and denominator using the *iterative* Euclid’s algorithm (see Algorithm 2 below).
7. Write a function, `reduce_fraction_inplace()`, to reduce a fraction in place using either of the `gcd()` functions. Add a short comment to your source code that states which of the two `gcd()` functions your code is using (state the line numbers), and how you know it is using that one. Add the function declaration to `fraction_toolbox.hpp`.
8. Write a function, `add_fractions()`, that takes two `fraction` structures as input, and returns their sum as a new `fraction` in it’s most reduced form.
9. Write a function, `sum_fraction_array_approx()`, that takes a pointer to an array of fractions and `n` (the length of the array) as input, and returns a double precision floating point number approximately equal to the sum of the array. Please write a brief comment in the code that explains why this is an approximate function.
10. Write a function, `sum_fraction_array()`, that takes a pointer to an array of fractions and `n` (the length of the array) as input, and returns a `fraction` equal to the sum of the array, by summing the `fraction` objects.

Euclid's Algorithm

[Euclid's Algorithm](#), discovered by the Greek mathematician Euclid in 300BC, is a method for finding the greatest common divisor of two integers. The algorithm can be implemented either [recursively](#) or iteratively, below we provide the pseudocode for both versions.

Please note: The \leftarrow symbol represents the assignment operator. In most programming languages this operator is represented by the $=$ symbol, although mathematically speaking this choice of symbols is questionable.

Algorithm 1: Euclidean Algorithm - Recursive

Data: Two integers a and b

Result: Integer greatest common divisor of a and b

```
1 Function gcd( $a, b$ ):  
2   if  $b = 0$  then  
3     return  $a$   
4   else  
5     return gcd( $b, a \bmod b$ )  
6   end
```

Algorithm 2: Euclidean Algorithm - Iterative

Data: Two integers a and b

Result: Integer greatest common divisor of a and b

```
1 Function gcd( $a, b$ ):  
2   while  $b \neq 0$  do  
3      $t \leftarrow b$   
4      $b \leftarrow a \bmod b$   
5      $a \leftarrow t$   
6   end  
7 return  $a$ 
```

3 Testing the Function Toolbox [30 points]

The `main.cpp` file is where you will find the skeleton of the testing program for your fraction toolbox. All of the functions that you declared in the header file will be available in your main program due to the include statement:

```
#include "fraction_toolbox.hpp"
```

Since compiling this program requires managing dependencies across different files that must be compiled, we use a makefile to automate compiling. To compile code using a makefile simply use:

```
[user@icsnodeXX]$ make
```

You may also clean out all compiled binaries by using:

```
[user@icsnodeXX]$ make clean
```

The `main.cpp` file also contains the `main()` function. In C code, this is the only function that will be executed when the code is run, thus if you want other functions to be executed when you run the code, you must call them in the `main()` function. Keep this in mind as you are building and testing the various other parts of this assignment.

Please implement the following functions in the `main.cpp` file to test your toolbox:

1. Write a function, `test23467()`, that reads command line arguments (you may use the pre-implemented `readcmdline()` function) to create a fraction, then tests functions 2, 3, 4, 6, and 7 by calling the functions and printing the results. Please also print descriptions of all of the results so that it is obvious which results represent which function. There is a pre-implemented print function for fractions, you can write your own print statements for testing functions 4 and 6 [5 points].

2. Write a function, `test5()`, that prompts the user for two integers and then prints the greatest common divisor of those two integers. **Hint:** check out the `cin` object [5 points].
3. Write a function, `test_array_functions()`, that:
 - (a) takes integer n as an input;
 - (b) dynamically allocates a memory buffer for an array of fractions of length n using the `malloc` function;
 - (c) passes a pointer to that memory buffer to the pre-implemented function `fill_fraction_array()`. The buffer will then be filled with an array of fractions of the form $\left[\frac{1}{1(1+1)}, \frac{1}{2(1+2)}, \dots, \frac{1}{n(1+n)} \right]$.
Hint: see pre-implemented `print_fraction_array()` function to verify results;
 - (d) passes a pointer to the (now filled) memory buffer to your `sum_fraction_array()` function;
 - (e) prints resultant sum;
 - (f) passes a pointer to the memory buffer to your `sum_fraction_array_approx()` function;
 - (g) prints the resultant sum.

Test this function in the `main()` function and ensure that it is giving you the results you would expect for small n . At some large n the `sum_fraction_array()` function will start giving funny results. Find that n , and then write a brief explanation of what is happening as a comment in the source code. Why is this not happening to the approximate function?
Hint: the values of 32-bit integers are between -2^{31} and $+2^{31}$. [20 points].
4. Once you have finished implementing the testing statements please call the `test_toolbox()` function from the `main()` function to prepare the code for submission.

Additional notes and submission details

Submit the source code files (together with your used `Makefile`) in an archive file (tar, zip, etc.) to [iCorsi](#).

- Your submission should be a gzipped tar archive, formatted like `project_number_lastname_firstname.zip` or `project_number_lastname_firstname.tgz`. It should contain:
 - the screenshot of your terminal for exercise 1.
 - all the source codes of your solutions.
- Submit your .zip or .tgz through [iCorsi](#).

Code of Conduct and Policy

- Do not use or otherwise access any on-line source or service other than the iCorsi system for your submission. In particular, you may not consult sites such as GitHub Co-Pilot or ChatGPT.
- You must acknowledge any code you obtain from any source, including examples in the documentation or course material. Use code comments to acknowledge sources.
- Your code must compile with a standard-configuration C/C++ compiler.